

Part 2:

Supervised vs Unsupervised Learning

Supervised Learning

- We have **input data X** and **output labels y** .
- Learn a function f such that $y \approx f(X)$.
- Examples: predicting house prices, classifying emails as spam / not spam.

Unsupervised Learning

- Only **input data X** (no labels).
- Discover structure or patterns in the data.
- Examples: customer segmentation, clustering images, PCA.

Two Ways to Learn from Data



Supervised Learning



Learns from data with answers.



Unsupervised Learning

Finds hidden patterns in data without answers.

800 × 401

Machine Learning
has two main
learning styles :

**just like students
learn in two
different ways**

Supervised Learning

Definition

- Computer learns using data **with answers**
- Inputs + Outputs available

Simple idea:

Learning with a teacher

Unsupervised Learning

Definition

- Computer learns using data **without answers**
- Only inputs available

Simple idea:

Learning without a teacher



Supervised Learning Example

Studying Past Exam Papers

You have:

- Questions
- Answers

You learn to solve new questions.

Real-Life Uses:

- Predict house prices
- Detect spam emails
- Predict student marks



Unsupervised Learning Example

Friend Groups in Classroom

No one tells the groups.

But you observe and find:

- Gamers
- Studious students
- Sports lovers

Real-Life Uses:

- Customer segmentation
- Grouping photos
- Market analysis

Examples of Machine Learning Algorithms

Supervised Learning

- Linear Regression
- Logistic Regression
- k -Nearest Neighbors (k-NN)
- Decision Trees
- Random Forests
- Support Vector Machines (SVM)
- Neural Networks

Unsupervised Learning

- k -Means Clustering
- Hierarchical Clustering
- DBSCAN
- Principal Component Analysis (PCA)
- Independent Component Analysis (ICA)
- Autoencoders

Python for Machine Learning

Python is widely used in ML because:

- Clean and readable syntax.
- Rich scientific libraries (NumPy, Pandas, Matplotlib).
- Strong ML ecosystem (scikit-learn, TensorFlow, PyTorch).
- Large community and good documentation.

Typical ML workflow in Python:

- Load and clean data.
- Explore and visualize.
- Build and evaluate models.
- Deploy and use models.

```
coffee.py

def make_coffee(*, with_sugar=False, with_milk=False):
    pass

make_coffee(with_milk=True, with_sugar=True)
```

All important Machine
[Machine Learning Functions](#)



Important Python Libraries for ML

Data Handling

- **NumPy** – arrays, matrices, math.
- **Pandas** – tables, CSV, Excel.
- **SciPy** – scientific routines.

ML + Visualization

- **Matplotlib** – plotting graphs.
- **scikit-learn** – classical ML algorithms.
- **Seaborn** – statistical plots.

Eg: open csv file **without** using Libraries

```
def read_csv_without_libraries(filename):  
    data = []  
    with open(filename, 'r') as file:  
        for line in file:  
            values = line.strip().split(',')  
  
            data.append(values)  
    return data  
csv_data = read_csv_without_libraries('example.csv')
```

Eg: open csv file **using** Pandas Library

```
import pandas as pd  
  
df = pd.read_csv('example.csv')
```

Basic Python Syntax for ML

Variables

```
a = 10, name = "USJ"
```

Lists and Arrays

- Python list: `x = [1, 2, 3]`
- NumPy array: `x = np.array([1, 2, 3])`

Function

```
def square(x):  
    return x * x
```

Importing Packages

```
import numpy as np import pandas as pd
```

Code:

basic python

https://github.com/LakshithaSenavirathna/MachineLearningTutorial/blob/main/ML_tutorial_1.ipynb

D A T A S E T



Datasets: Features, Labels, Train/Test Split

Features and Labels

- **Features** (inputs): variables used to predict.
- **Label** (target): what we want to predict.
- In code: features X , labels y .

Train / Test Split

- Split the data into:
 - **Training set** – used to fit the model.
 - **Test set** – used to evaluate performance.
- Helps estimate generalization to unseen data.



scikit-learn Interface Pattern

Most models in `scikit-learn` follow this pattern:

- ① **Create** a model `model = LinearRegression()`
- ② **Fit** on training data `model.fit(X_train, y_train)`
- ③ **Predict** on new data `y_pred = model.predict(X_test)`

Same interface for many algorithms: trees, SVM, k-NN, etc.



```
import pandas as pd

# Load the CSV file
data = pd.read_csv (" simple_regression_data .csv ")

# Look at the first few rows
print ( data . head ())

# Summary statistics
print ( data . describe ())
```

Code:

load csv

https://github.com/LakshithaSenavirathna/MachineLearningTutorial/blob/main/ML_tutorial_2.ipynb

Simple Linear Regression with One Variable

Model

Assume a linear relationship between x and y :

$$y = \beta_0 + \beta_1 x + \varepsilon,$$

where:


- β_0 – intercept,
- β_1 – slope,
- ε – error term.

Idea

- Observe data points (x_i, y_i) .
- Choose β_0, β_1 so the line is “close” to the points.
- Minimize the sum of squared errors:

$$SSE = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Visualizing Simple Linear Regression



Requirements for the Data Set

- Numerical variables: one predictor x and one response y .
- Enough points (e.g. $n \geq 20$) to see the relationship.
- Approximate linearity in the scatter plot of y vs x .
- No extreme outliers dominating the fit.
- Reasonable spread in x ; not all x_i almost equal.

Exploring and Visualizing the Data

Statistical Summaries

- Mean, median, standard deviation of x and y .
- Minimum, maximum, quartiles.
- Correlation between x and y .

Graphics

- Scatter plot of y vs x .
- Histogram of residuals after fitting.
- Residual vs fitted plot to check assumptions.

Code

https://github.com/LakshithaSenavirathna/MachineLearningTutorial/blob/main/ML_tutorial_3.ipynb

scikit-learn Pipeline (Conceptual)

A typical scikit-learn pipeline:

- **Step 1: Preprocessing**
 - Handle missing values.
 - Scale / normalize features.
 - Encode categorical variables.
- **Step 2: Model**
 - LinearRegression, RandomForestRegressor, etc.
- **Step 3: Prediction**
 - Use the trained pipeline to predict on new data.

In code (idea): `Pipeline([("scaler", StandardScaler()), ("model", LinearRegression())])`

Overall Python ML Workflow

- ① Collect data.
- ② Clean and preprocess data.
- ③ Explore and visualize data.
- ④ Choose and train ML model(s).
- ⑤ Evaluate on test data.
- ⑥ Deploy the model and monitor performance.

Each step can be demonstrated with short Python examples in a Jupyter notebook.

Thank You 😊